

# Variational Autoencoders (VAEs)

Roko Čubrić

September 2025

## 1 Introduction to Standard Autoencoders

A standard **autoencoder** is a neural network model used for unsupervised learning of data representations. It consists of two main parts:

1. An **encoder**, which compresses a high-dimensional input  $x$  into a low-dimensional latent representation  $z$ .
2. A **decoder**, which takes the latent representation  $z$  and attempts to reconstruct the original input  $\hat{x}$ .

The objective of a standard autoencoder is to minimize the reconstruction loss, which is typically the mean squared error (MSE) or binary cross-entropy between the original input  $x$  and the reconstructed output  $\hat{x}$ . This forces the model to learn a compact representation that still retains enough information to reproduce the input.

However, standard autoencoders have a significant limitation: the learned latent space is not guaranteed to be **continuous** or **structured**. They may simply map each input to an isolated, non-overlapping point. When we try to generate new data by sampling a random point from this latent space, the decoder often produces nonsensical output because the sampled point falls into a "gap" where the model has learned nothing.

## 2 The Variational Autoencoder (VAE)

The **Variational Autoencoder (VAE)** addresses this limitation by taking a fundamentally different approach. It is a generative model that aims to learn not just a point representation, but a **probability distribution** for the latent space. This ensures the latent space is continuous and structured, allowing us to generate new, meaningful data.

The VAE's objective is to maximize the **log-likelihood** of the data, which means we want to maximize the probability  $p_\theta(x)$  of our input data (something resembling it) being generated by our model. A higher probability indicates our model has learned to represent the data well. This probability is expressed as:

$$\log p_\theta(x) = \log \int p_\theta(x, z) dz$$

This integral is computationally intractable (not theoretically solvable). To solve this, the VAE uses a technique called **variational inference**, where we introduce a simple, tractable distribution  $q_\phi(z|x)$  to approximate the true but intractable posterior  $p_\theta(z|x)$ . The core idea is that we can find a lower bound on our objective, which is easier to optimize. This lower bound is the **Evidence Lower Bound (ELBO)**.

### Example: Why Use Log-Likelihood?

When working with probabilities in machine learning, we often multiply many very small numbers together. For example, the probability of a specific image appearing is a product of the probabilities of all its individual pixels. If an image has  $D$  pixels, the total probability is a product of  $D$  small numbers.

$$p(x) = p(x_1) \cdot p(x_2) \cdot \dots \cdot p(x_D)$$

Since these probabilities are between 0 and 1, multiplying thousands or millions of them together would quickly lead to a number so close to zero that a computer cannot represent it. This is called a **floating-point underflow**, which effectively makes the probability zero and breaks the training process.

To avoid this, we take the logarithm of the probability. A key property of logarithms is that they convert multiplication into summation:

$$\log(p(x)) = \log(p(x_1)) + \log(p(x_2)) + \dots + \log(p(x_D))$$

Summing up numbers is far more numerically stable for a computer than multiplying them. By maximizing the log-likelihood (the sum), we achieve the same goal as maximizing the original probability, but without the risk of numerical errors.

## 2.1 Deriving the ELBO

The **Evidence Lower Bound (ELBO)** is the lower bound on the log-likelihood that we aim to optimize. It can be derived as follows:

$$\log p_\theta(x) = \log \int p_\theta(x, z) dz = \log \int \frac{p_\theta(x, z)}{q_\phi(z|x)} q_\phi(z|x) dz$$

**Lemma 1** (Jensen's Inequality). *For a concave function  $f$  and a random variable  $X$ , the following inequality holds:  $f(\mathbb{E}[X]) \geq \mathbb{E}[f(X)]$ . Since the logarithm function is concave, we can use this inequality to find a lower bound.*

Applying Jensen's Inequality to the equation above, we get the lower bound:

$$\log p_\theta(x) \geq \int q_\phi(z|x) \log \left( \frac{p_\theta(x, z)}{q_\phi(z|x)} \right) dz$$

This expectation can be broken down using the properties of logarithms and probability distributions.

$$\begin{aligned} & \int q_\phi(z|x) [\log p_\theta(x, z) - \log q_\phi(z|x)] dz \\ & \int q_\phi(z|x) [\log p_\theta(x|z) + \log p(z) - \log q_\phi(z|x)] dz \end{aligned}$$

Rearranging and grouping terms, we get the final form of the ELBO. First, we separate the integral into two terms:

$$\int q_\phi(z|x) \log p_\theta(x|z) dz - \int q_\phi(z|x) [\log q_\phi(z|x) - \log p(z)] dz$$

Recognizing the definitions of an expectation and the Kullback-Leibler (KL) divergence, we can rewrite the terms as:

$$\mathbb{E}_{z \sim q_\phi} [\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x) \parallel p(z))$$

This gives us the final form of the ELBO:

$$\text{ELBO}(x) = \mathbb{E}_{z \sim q_\phi} [\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x) \parallel p(z))$$

This formula is at the heart of the VAE. It is a sum of two terms that are optimized simultaneously. The loss function used for training is the negative of the ELBO.

## 2.2 The Two Terms of the VAE Loss Function

The ELBO's two terms have distinct roles and are calculated for each specific data point  $x$  during training.

### 2.2.1 The Reconstruction Term

$$\mathbb{E}_{z \sim q_\phi}[\log p_\theta(x|z)] = \int q_\phi(z|x) \log p_\theta(x|z) dz$$

This term measures how well the decoder can reconstruct the original data point  $x$  from the latent codes  $z$  provided by the encoder. It is a log-likelihood, and its maximization corresponds to minimizing a reconstruction loss (e.g., binary cross-entropy for images). This term trains both the **encoder** and the **decoder**. Unlike a standard autoencoder that predicts a single pixel value, the VAE's decoder outputs the parameters of a probability distribution for each pixel. The final image is created by taking the mean values, but the variance is a crucial component of the loss function. It allows the model to express uncertainty about its predictions; a larger variance will reduce the reconstruction penalty if a prediction is slightly off, which helps the model learn a more robust representation. This is the only part of the VAE that uses the variance of the pixel distributions.

#### Example: Log-Likelihood Formulas

The formula for  $\log p_\theta(x|z)$  depends on the type of data.

For continuous data (e.g., grayscale images with pixel values from 0 to 1), the decoder outputs a Gaussian distribution for each pixel. The log-likelihood is:

$$\log p_\theta(x|z) = \sum_{i=1}^D \log \mathcal{N}(x_i; \mu_{\theta,i}(z), \sigma_{\theta,i}^2(z))$$

Here,  $D$  is the dimensionality of the data, which is the total number of pixels in the image. The total loss is the sum of the negative log-likelihoods for all pixels, where the loss for a single pixel is:

$$-\log \mathcal{N}(x_i; \mu_{\theta,i}, \sigma_{\theta,i}^2) = \frac{1}{2} \log(2\pi) + \frac{1}{2} \log(\sigma_{\theta,i}^2) + \frac{(x_i - \mu_{\theta,i})^2}{2\sigma_{\theta,i}^2}$$

For binary data (e.g., black and white images with pixel values of 0 or 1), the decoder outputs a Bernoulli distribution for each pixel. The log-likelihood is:

$$\log p_\theta(x|z) = \sum_{i=1}^D [x_i \log(\hat{x}_i) + (1 - x_i) \log(1 - \hat{x}_i)]$$

where  $\hat{x}_i$  is the decoder's predicted probability for the  $i$ -th pixel, and this is equivalent to the negative of the binary cross-entropy loss.

### 2.2.2 The KL Divergence Term

$$D_{KL}(q_\phi(z|x) \| p(z)) = \int q_\phi(z|x) \log \frac{q_\phi(z|x)}{p(z)} dz$$

This is a regularization term that measures the difference between the encoder's output distribution,  $q_\phi(z|x)$ , and a predefined prior distribution,  $p(z)$  (typically a standard normal distribution,  $\mathcal{N}(0, I)$ ). By minimizing this term, we force the latent distributions for all our data points to be organized and centered around the origin, which creates a continuous and structured latent space. This term trains only the **encoder**. The analogy is that instead of isolated data points on a map, this term forces the data points to be on well-defined streets that can be smoothly navigated.

## 3 Backpropagation and the Reparameterization Trick

Since the VAE trains using backpropagation, we need a way to pass gradients from the loss function all the way back to the encoder's parameters. A critical problem arises when we try to sample a random

latent variable  $z$  from our distribution  $q_\phi(z|x)$ . The act of random sampling is a non-differentiable operation, which would break the backpropagation chain.

The **reparameterization trick** solves this problem by separating the randomness from the model's parameters. Instead of sampling from  $q_\phi(z|x)$ , we sample a random variable  $\epsilon$  from a simple standard normal distribution,  $\mathcal{N}(0, I)$ , and then compute  $z$  as a deterministic, differentiable function of  $\mu$  and  $\sigma$  (the mean and standard deviation output by the encoder).

$$z = \mu + \sigma \cdot \epsilon$$

This allows us to backpropagate gradients from the reconstruction loss back through the encoder.

- The gradient with respect to  $\mu$  is straightforward:  $\frac{\partial \mathcal{L}}{\partial \mu} = \frac{\partial \mathcal{L}}{\partial z} \cdot \frac{\partial z}{\partial \mu} = \frac{\partial \mathcal{L}}{\partial z} \cdot 1$ .
- The gradient with respect to  $\sigma$  is now also possible:  $\frac{\partial \mathcal{L}}{\partial \sigma} = \frac{\partial \mathcal{L}}{\partial z} \cdot \frac{\partial z}{\partial \sigma} = \frac{\partial \mathcal{L}}{\partial z} \cdot \epsilon$ .

The trick is essential for the variance part of the backpropagation as it makes the random sampling step differentiable.

The reparameterization trick is more accurately viewed not as a trick, but as the only logical method for computing the latent variable  $z$  in a way that allows for the necessary backpropagation and gradient calculation.

## 4 Amortized Inference and Latent Space Arithmetic

The VAE's design allows for other key technical advantages.

### Amortized Inference

During training, we "amortize" the cost of finding the approximate posterior. Instead of running a separate optimization for each data point, we train a single neural network (the encoder) to learn a general mapping from any data point to its latent distribution parameters. This makes the VAE incredibly efficient at inference time, as it can instantly find a latent representation for any new input. This is analogous to doing all the prep work for a feast (e.g., chopping all vegetables) once, so that each dish can be cooked quickly.

### Latent Space Arithmetic

Because the VAE learns a continuous and structured latent space, we can perform meaningful arithmetic on latent vectors. A famous example is manipulating the concept of "glasses." We can encode an image of a person with glasses ( $z_{\text{with glasses}}$ ) and an image of the same person without glasses ( $z_{\text{without glasses}}$ ), and then find a vector representing the "glasses" concept:

$$z_{\text{glasses}} = z_{\text{with glasses}} - z_{\text{without glasses}}$$

This "glasses" vector can then be added to the latent vector of a different person, and the decoder can produce an image of that person with glasses. This is a direct result of the structured latent space created by the KL divergence term.

## 5 Conclusion

The VAE is a powerful generative model that overcomes the limitations of a standard autoencoder by learning a continuous and structured latent space. This is achieved through the ELBO, a unique loss function that balances reconstruction quality with regularization. The reparameterization trick allows for efficient training via backpropagation, and concepts like amortized inference enable the model to be practical for a wide range of applications.